

egefexilu

April 29, 2024

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
from sklearn.feature_selection import SelectKBest
from sklearn.ensemble import RandomForestClassifier
from sklearn.pipeline import make_pipeline
from sklearn.model_selection import cross_validate
from collections import Counter
from xgboost import XGBClassifier
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier,
↳ GradientBoostingClassifier, ExtraTreesClassifier, VotingClassifier
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV, cross_val_score,
↳ StratifiedKFold, learning_curve
```

1 Introduction

- Acquiring Data and analyzing the dataset
- EDA - Exploratory data analysis

2 Acquiring Data

We'll first read the csv(comma separated values) files into a training and testing dataframes which we can manipulate for analysis and predictions

```
[ ]: #Reading data in the .csv format (no special encoding schemes required)

df_train = pd.read_csv('/content/train.csv')
```

```
df_test = pd.read_csv('/content/test.csv')
id_t = df_test['PassengerId']
print(df_train.shape)
print(df_test.shape)
```

```
(891, 12)
```

```
(418, 11)
```

Outlier detection

```
[ ]: #Outlier detection using the Tukey method

def detect(dataframe,n,features):
    ind = []
    for column in features:
        q1 = np.percentile(dataframe[column],25)
        q3 = np.percentile(dataframe[column],75)
        res = q3 - q1
        new = 1.5 * res
        out = dataframe[(dataframe[column]<q1 - new) | (dataframe[column] >q3 +
↪new)].index
        ind.extend(out)
    ind = Counter(ind)
    val = list(s for s,i in ind.items() if i>n)
    return val

#Detecting outliers for the numerical features
ott = detect(df_train,2,["Age","SibSp","Parch","Fare"])
```

Outliers can have dramatic effects on our predictions and can also result in harming our final result. We find outliers for the numerical features in the given dataset and we outlined rows with atleast two outliers

```
[ ]: #Our final outliers - can have a harmful effect on regression tasks
df_train.loc[ott]
```

```
[ ]:
   PassengerId  Survived  Pclass                                Name  Sex \
27           28         0        1  Fortune, Mr. Charles Alexander   male
88           89         1        1    Fortune, Miss. Mabel Helen  female
159          160         0        3    Sage, Master. Thomas Henry   male
180          181         0        3  Sage, Miss. Constance Gladys  female
201          202         0        3      Sage, Mr. Frederick       male
324          325         0        3    Sage, Mr. George John Jr   male
341          342         1        1  Fortune, Miss. Alice Elizabeth  female
792          793         0        3      Sage, Miss. Stella Anna  female
846          847         0        3      Sage, Mr. Douglas Bullen   male
863          864         0        3  Sage, Miss. Dorothy Edith "Dolly"  female
```

	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
27	19.0	3	2	19950	263.00	C23 C25 C27	S
88	23.0	3	2	19950	263.00	C23 C25 C27	S
159	NaN	8	2	CA. 2343	69.55	NaN	S
180	NaN	8	2	CA. 2343	69.55	NaN	S
201	NaN	8	2	CA. 2343	69.55	NaN	S
324	NaN	8	2	CA. 2343	69.55	NaN	S
341	24.0	3	2	19950	263.00	C23 C25 C27	S
792	NaN	8	2	CA. 2343	69.55	NaN	S
846	NaN	8	2	CA. 2343	69.55	NaN	S
863	NaN	8	2	CA. 2343	69.55	NaN	S

These rows were identified as outliers according to our given function so we will proceed to drop them from our dataframe

```
[ ]: df_train.drop(ott,axis=0,inplace=True) #dropping on the row axis
```

Create a final dataset (used while feature tuning and cleaning) by combining the training and testing dataframe

```
[ ]: #Dataset combining both the training and testing datasets
t_len = len(df_train)
final = pd.concat([df_train,df_test],axis=0).reset_index(drop=True)
final.shape
```

```
[ ]: (1299, 12)
```

3 Analyzing the dataset

Viewing the first 5 rows of the dataset

```
[ ]: df_train.head()
```

```
[ ]: PassengerId  Survived  Pclass  \
0             1         0         3
1             2         1         1
2             3         1         3
3             4         1         1
4             5         0         3
```

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

List of columns present in the dataset

```
[ ]: df_train.columns
```

```
[ ]: Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',
          'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked'],
          dtype='object')
```

```
[ ]: df_test.head()
```

```
[ ]:
   PassengerId  Pclass                               Name  Sex \
0          892      3                        Kelly, Mr. James  male
1          893      3      Wilkes, Mrs. James (Ellen Needs)  female
2          894      2                Myles, Mr. Thomas Francis  male
3          895      3                Wirz, Mr. Albert  male
4          896      3  Hirvonen, Mrs. Alexander (Helga E Lindqvist)  female
```

	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	34.5	0	0	330911	7.8292	NaN	Q
1	47.0	1	0	363272	7.0000	NaN	S
2	62.0	0	0	240276	9.6875	NaN	Q
3	27.0	0	0	315154	8.6625	NaN	S
4	22.0	1	1	3101298	12.2875	NaN	S

Checking for NULL values

Multiple NULL values present in Age and Cabin while only 2 in the Embarked column.

```
[ ]: df_train.info()
      df_train.isnull().sum()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 881 entries, 0 to 890
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PassengerId  881 non-null    int64
1   Survived     881 non-null    int64
2   Pclass       881 non-null    int64
3   Name         881 non-null    object
4   Sex          881 non-null    object
5   Age          711 non-null    float64
```

```

6  SibSp      881 non-null  int64
7  Parch      881 non-null  int64
8  Ticket     881 non-null  object
9  Fare       881 non-null  float64
10 Cabin      201 non-null  object
11 Embarked   879 non-null  object
dtypes: float64(2), int64(5), object(5)
memory usage: 89.5+ KB

```

```

[ ]: PassengerId      0
     Survived         0
     Pclass           0
     Name             0
     Sex              0
     Age             170
     SibSp            0
     Parch            0
     Ticket           0
     Fare             0
     Cabin           680
     Embarked         2
     dtype: int64

```

```

[ ]: df_test.info()
     df_test.isnull().sum()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 418 entries, 0 to 417
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     418 non-null   int64
1   Pclass          418 non-null   int64
2   Name            418 non-null   object
3   Sex             418 non-null   object
4   Age             332 non-null   float64
5   SibSp           418 non-null   int64
6   Parch           418 non-null   int64
7   Ticket          418 non-null   object
8   Fare            417 non-null   float64
9   Cabin           91 non-null    object
10  Embarked        418 non-null   object
dtypes: float64(2), int64(4), object(5)
memory usage: 36.0+ KB

```

```

[ ]: PassengerId      0
     Pclass           0
     Name             0

```

```

Sex            0
Age            86
SibSp          0
Parch          0
Ticket         0
Fare           1
Cabin         327
Embarked       0
dtype: int64

```

The describe method provides us a brief summary of basic statistical measures such as quartiles, mean, count etc.

```
[ ]: #Brief statistics
df_train.describe()
```

```
[ ]:
      PassengerId  Survived  Pclass    Age  SibSp  \
count    881.000000    881.000000    881.000000  711.000000  881.000000
mean      446.713961     0.385925     2.307605   29.731603    0.455165
std       256.617021     0.487090     0.835055   14.547835    0.871571
min         1.000000     0.000000     1.000000     0.420000    0.000000
25%       226.000000     0.000000     2.000000   20.250000    0.000000
50%       448.000000     0.000000     3.000000   28.000000    0.000000
75%       668.000000     1.000000     3.000000   38.000000    1.000000
max       891.000000     1.000000     3.000000   80.000000    5.000000

      Parch    Fare
count    881.000000  881.000000
mean       0.363224   31.121566
std       0.791839   47.996249
min        0.000000    0.000000
25%        0.000000    7.895800
50%        0.000000   14.454200
75%        0.000000   30.500000
max         6.000000  512.329200

```

```
[ ]: df_test.describe()
```

```
[ ]:
      PassengerId  Pclass    Age  SibSp  Parch    Fare
count    418.000000  418.000000  332.000000  418.000000  418.000000  417.000000
mean    1100.500000   2.265550  30.272590   0.447368   0.392344   35.627188
std     120.810458   0.841838  14.181209   0.896760   0.981429   55.907576
min      892.000000   1.000000   0.170000   0.000000   0.000000    0.000000
25%     996.250000   1.000000  21.000000   0.000000   0.000000    7.895800
50%    1100.500000   3.000000  27.000000   0.000000   0.000000   14.454200
75%    1204.750000   3.000000  39.000000   1.000000   0.000000   31.500000
max    1309.000000   3.000000  76.000000   8.000000   9.000000  512.329200


```

4 EDA (exploratory data analysis)

Checking the correlation among the numerical variables present in the dataset

```
[ ]: #Correlation between numerical values
```

```
hm = sns.heatmap(df_train[['Survived','SibSp','Parch','Age','Fare']].  
    ↪corr(),annot=True,cmap='Blues')
```



We see that there's high correlation only between fare and survival probabilities, this feature may turn out to be our main focus


```
[ ]: df_train[['Sex','Survived']].groupby(['Sex'],as_index=False).mean().  
    ↪sort_values(by='Survived',ascending=False)
```

```
[ ]:      Sex  Survived  
0  female  0.747573  
1   male   0.190559
```

There is a much higher probability of females surviving on the Titanic than men, Let us now proceed to visualize it.

```
[ ]: #Plot for gender distribution of survival rate
sns.barplot(x='Sex',y='Survived',data=df_train)
plt.title('Survival %: Gender distribution')
plt.ylabel('% of passengers survived')
plt.xlabel('Sex')
```

```
[ ]: Text(0.5, 0, 'Sex')
```




We can clearly see that gender plays a huge role in calculating the survival probability for a passenger. This is a crucial feature that would help to improve upon our predictions.

We will now visualize the total passenger count according to gender and survivability.

```
[ ]: #Survivor count for each gender
m = sns.countplot(x='Sex',hue='Survived',data=df_train)
plt.title('Survival count: Gender distribution')
plt.xlabel('Sex')
plt.ylabel('Number of passengers survived')
a = m.get_legend()
b = a.texts
b[0].set_text('No')
```




```
b[1].set_text('Yes')
```



Again, we see that there is a large difference between the number of survivors between both of the genders.

```
[ ]: #Age survival ratio  
s = sns.FacetGrid(df_train,col='Survived')  
s.map(plt.hist,'Age',bins=20)
```

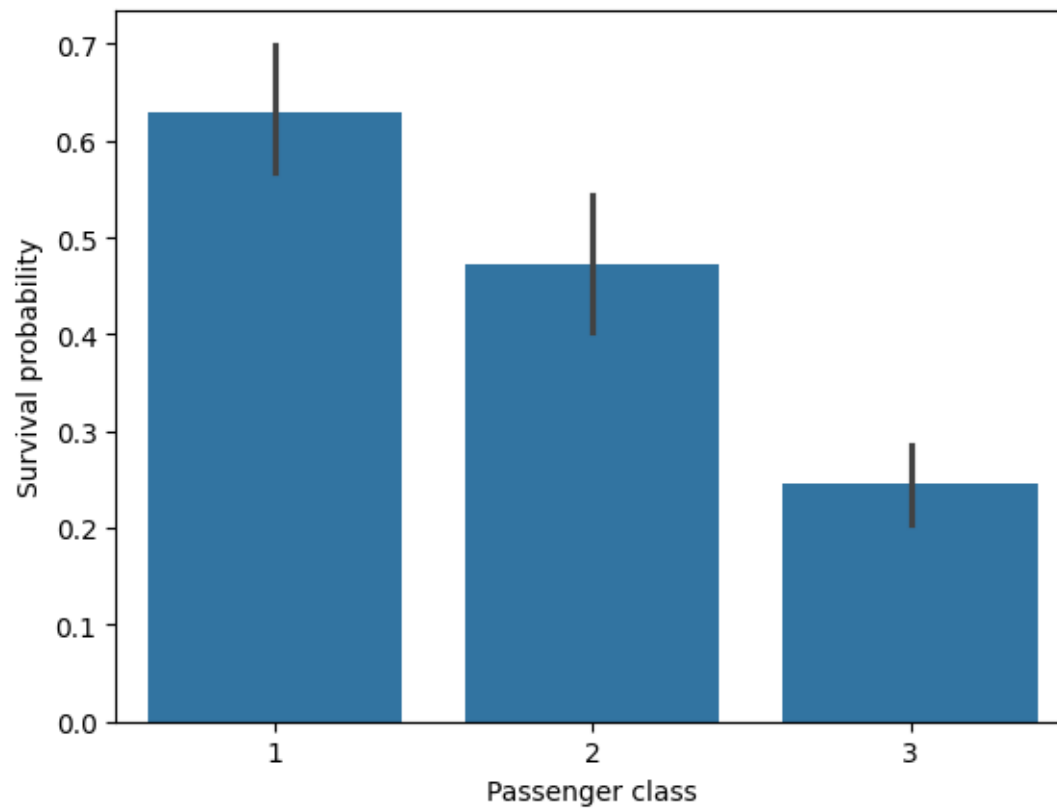
```
[ ]: <seaborn.axisgrid.FacetGrid at 0x7bd1d187b3a0>
```



It seems that the middle-aged and old people had a very low survival probability as compared to infants, childrens and young adults. Age plays an important role in saving yourself during a time of crisis.


```
[ ]: #Survival probability vs Class
b = sns.barplot(x='Pclass',y='Survived',data=df_train)
b.set_ylabel('Survival probability')
b.set_xlabel('Passenger class')
```

```
[ ]: Text(0.5, 0, 'Passenger class')
```



```
[ ]: c = sns.barplot(x='Pclass',y='Fare',data=df_train)
c.set_ylabel('Fare')
c.set_xlabel('Pclass')
```

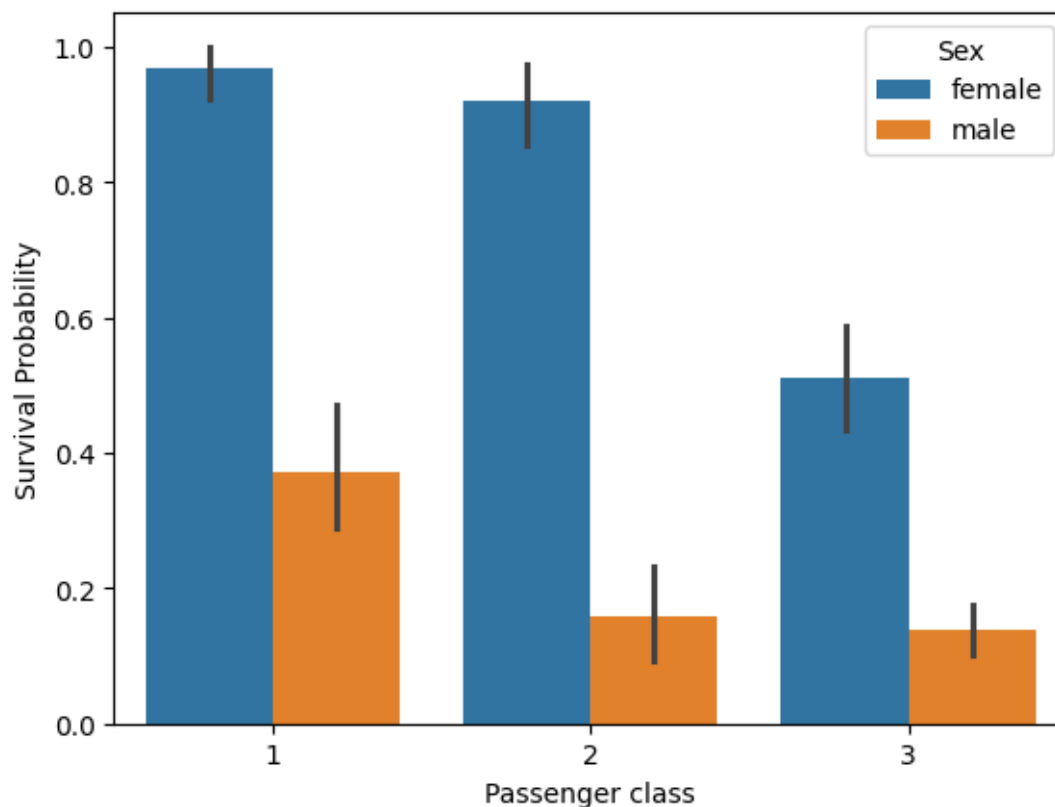
```
[ ]: Text(0.5, 0, 'Pclass')
```



It seems that the rich people present in Passenger Class 1 had a much higher probability of survival in comparison to the other passenger classes. The fare is proportional to the survivability result.

```
[ ]: #Survival Probability vs Class and Sex  
b = sns.barplot(x='Pclass',y='Survived',hue='Sex',data=df_train)  
b.set_ylabel('Survival Probability')  
b.set_xlabel('Passenger class')
```


```
[ ]: Text(0.5, 0, 'Passenger class')
```



We observe a similar trend here where the survival probability is in direct correlation with the gender as well as the Passenger class. Higher the passenger class equalled low fare which results in low survival chance.

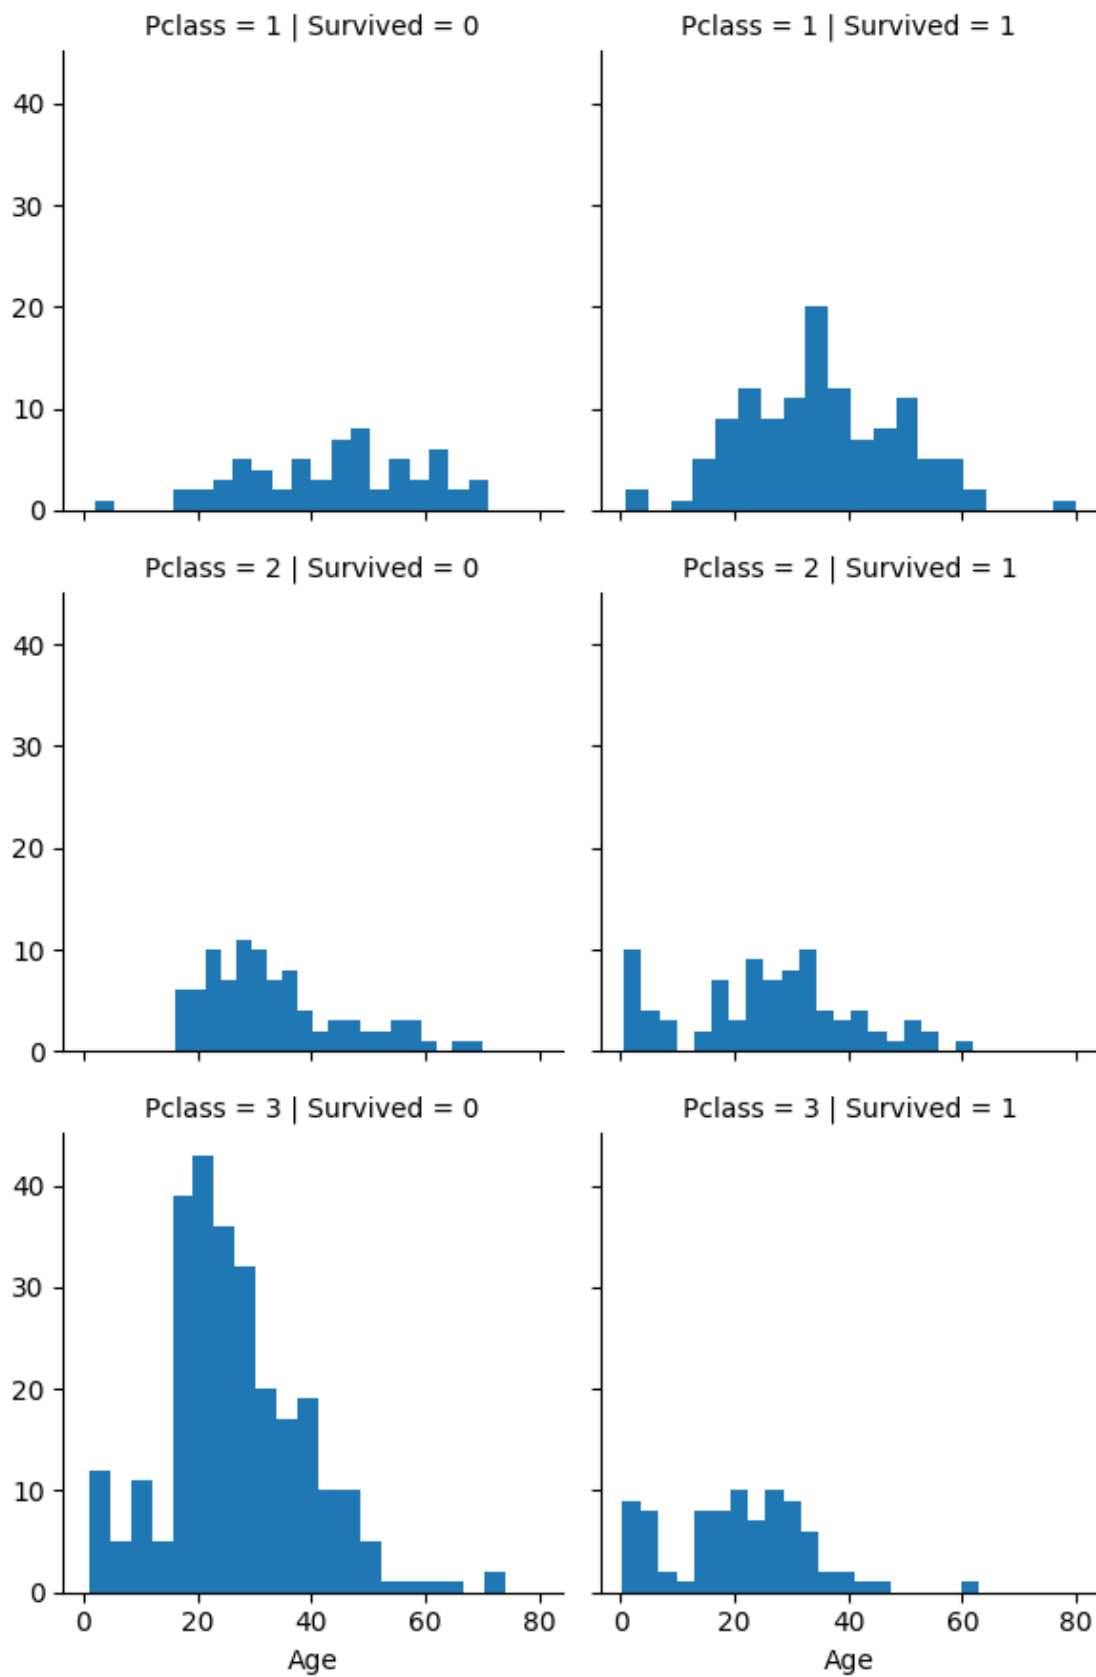
```
[ ]: #Exploring Embarked vs Survival probability
ax = sns.barplot(x='Embarked',y='Survived',data=df_train)
ax.set_ylabel('Survival probability')
ax.set_xlabel('Embarked')
```

```
[ ]: Text(0.5, 0, 'Embarked')
```



```
[ ]: g1 = sns.FacetGrid(df_train,col='Survived',row='Pclass')  
g1.map(plt.hist,'Age',bins=20)
```

```
[ ]: <seaborn.axisgrid.FacetGrid at 0x7bd1d16d6da0>
```




We see that Pclass 1 consists of more middle-aged and old men while P-class 2 and 3 consist of infants, children and the like. It is a similar occurrence in real-life where usually old and rich people travel in First-Class with the presence of children being very rare.

We also see that the survived histogram dwindles down in height as we move lower down the graph. The histogram still rises up pretty well for the lower age group which is what we saw above as well. Children, teenagers and young adults have a higher survival rate than the middle-aged and old people.


```
[ ]: #Fare distribution
a1 = sns.kdeplot(df_train.
    ↳loc[(df_train['Survived']==0), 'Fare'], shade=True, label='Not Survived')
a1 = sns.kdeplot(df_train.
    ↳loc[(df_train['Survived']==1), 'Fare'], shade=True, label='Survived')
plt.ylabel('Survival frequency')
plt.xlabel('Fare distribution')
```

```
[ ]: Text(0.5, 0, 'Fare distribution')
```




```
[ ]: #Survival by Fare, Age and Sex  
s = sns.FacetGrid(df_train, hue='Survived', col='Sex')  
s.map(plt.scatter, 'Fare', 'Age', edgecolor='w').add_legend()
```

```
[ ]: <seaborn.axisgrid.FacetGrid at 0x7bd1cef73970>
```



```
[ ]: #Plotting fare in the combined dataset  
s = sns.distplot(final['Fare'])
```

