

qbn3llhwk

April 29, 2024

```
[ ]: import pandas as pd
import numpy as np
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
import seaborn as sns
import matplotlib.pyplot as plt

dataset= pd.read_csv("diabetes_dataset.csv")
```

```
[ ]: dataset.describe()
```

```
[ ]:
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin \ |
|-------|-------------|------------|---------------|---------------|------------|
| count | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 |
| mean | 3.845052 | 120.894531 | 69.105469 | 20.536458 | 79.799479 |
| std | 3.369578 | 31.972618 | 19.355807 | 15.952218 | 115.244002 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 25% | 1.000000 | 99.000000 | 62.000000 | 0.000000 | 0.000000 |
| 50% | 3.000000 | 117.000000 | 72.000000 | 23.000000 | 30.500000 |
| 75% | 6.000000 | 140.250000 | 80.000000 | 32.000000 | 127.250000 |
| max | 17.000000 | 199.000000 | 122.000000 | 99.000000 | 846.000000 |

| | BMI | DiabetesPedigreeFunction | Age | Outcome |
|-------|------------|--------------------------|------------|------------|
| count | 768.000000 | 768.000000 | 768.000000 | 768.000000 |
| mean | 31.992578 | 0.471876 | 33.240885 | 0.348958 |
| std | 7.884160 | 0.331329 | 11.760232 | 0.476951 |
| min | 0.000000 | 0.078000 | 21.000000 | 0.000000 |
| 25% | 27.300000 | 0.243750 | 24.000000 | 0.000000 |
| 50% | 32.000000 | 0.372500 | 29.000000 | 0.000000 |
| 75% | 36.600000 | 0.626250 | 41.000000 | 1.000000 |
| max | 67.100000 | 2.420000 | 81.000000 | 1.000000 |

```
[ ]: from sklearn.model_selection import train_test_split
X= dataset.drop("Outcome", axis=1)
y= dataset[["Outcome"]]
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.30,
↳random_state=7)
```

```
[ ]: model = LogisticRegression()
model.fit(X_train, y_train)
y_predict = model.predict(X_test)
model_score= model.score (X_test, y_test)
model_score
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
 n_iter_i = _check_optimize_result(


```
[ ]: 0.7489177489177489
```

```
[ ]: from sklearn import metrics
cnf_matrix = metrics.confusion_matrix(y_test, y_predict)
print(cnf_matrix)
```

```
[[127  20]
 [ 38  46]]
```

```
[ ]: class_names=[0,1]
fig, ax = plt.subplots()
tick_marks = np.arange(len(class_names))
plt.xticks(tick_marks, class_names)
plt.yticks(tick_marks, class_names)
sns.heatmap(pd.DataFrame(cnf_matrix), annot=True, cmap="YlGnBu" ,fmt='g')
ax.xaxis.set_label_position("top")
plt.tight_layout()
plt.title('Confusion matrix', y=1.1)
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
```

```
[ ]: Text(0.5, 427.95555555555555, 'Predicted label')
```



```
[ ]: true_neg, false_pos, false_neg, true_pos = cnf_matrix.ravel()
true_neg, false_pos, false_neg, true_pos
total = true_neg + false_pos + false_neg + true_pos
```

```
[ ]: accuracy = (true_pos + true_neg)/total
print(accuracy)
```

0.7489177489177489

```
[ ]: precision = true_pos/(true_pos + false_pos)
print(precision)
```

0.696969696969697

```
[ ]: recall = true_pos/(true_pos + false_neg)
print(recall)
```

0.5476190476190477

```
[ ]: f1_score = (2*precision*recall)/(precision+recall)
      print(f1_score)
```

0.6133333333333334